

PRIYANSHU PARIHAR

priyanshuparihar207@gmail.com | +9837610363 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Aspiring AI/ML Engineer with hands-on experience building production-grade Retrieval-Augmented Generation (RAG) systems, LLM orchestration pipelines, and multi-database AI architectures (PostgreSQL, Qdrant, Neo4j). Skilled in Python, LangGraph, NLP, and AI safety/guardrail design, with direct experience deploying secure, low-latency conversational AI across healthcare and financial domains. Strong analytics foundation (SQL, Power BI) and fast-growing expertise in machine learning and generative AI engineering.

TECHNICAL SKILLS

Languages: Python, SQL, HTML, CSS
AI / Machine Learning: RAG, Self-RAG, LangGraph, LangSmith, Prompt Engineering, NLP, Agentic Workflows, Qwen3-Embedding, sentence-transformers.
Databases & Vector Stores: PostgreSQL, MongoDB, Qdrant, Neo4j (Knowledge Graph).
Backend & APIs: FastAPI, REST APIs
AI Safety & Governance: Guardrails AI, Llm-guard, PII Masking, Prompt-Injection Defense, Semantic Caching
Data Visualization & Analytics: Power BI, Matplotlib, Pandas, NumPy
Tools: Git, GitHub, Docker, VS Code

EXPERIENCE

Mobcoder — AI/ML Engineer Intern March 2026 – Present 2026

- Built and deployed two production-grade Retrieval-Augmented Generation (RAG) platforms spanning the financial and healthcare domains
- Worked across the AI stack — retrieval pipelines, agentic orchestration, vector/graph databases, and safety guardrails — to ship secure, low-latency conversational AI systems.

PROJECTS

Alphalens AI — Financial Document Intelligence & RAG Platform | Mobcoder Internship [Demo Link](#)

Tech Stack: Python, FastAPI, Next.js, Qdrant, LangGraph, Guardrails AI, sentence-transformers, Docling

- Built a hybrid RAG platform (dense vector + BM25 search with cross-encoder reranking) for financial PDF analysis, achieving sub-0.1s retrieval latency.
- Designed LangGraph-based agentic query routing and Self-RAG reflection loops to reduce hallucinations and improve answer quality.
- Added Guardrails AI for PII masking and prompt-injection defense, plus dual-layer caching delivering 0.04s repeat-query responses.

Medical-Grade AI Assistant & Knowledge Spine — Healthcare RAG | Mobcoder Internship [Demo Link](#)

Tech Stack: Python, PostgreSQL, Qdrant, Neo4j, LLaMA-3.3-70B, Phi-3.5, Qwen3-Embedding, Alembic, Llm-guard

- Architected a multi-database AI infrastructure (PostgreSQL, Qdrant, Neo4j) for a medical-grade conversational assistant with immutable audit logging.
- Built an LLM governance gateway with semantic caching and parallel Self-RAG retrieval for accurate, low-latency responses.
- Integrated Llm-guard adversarial scanning and an NLP-driven pharmacovigilance feature with citation-enforced safety guardrails.

EDUCATION

B.Tech, Information Technology — M.L.V. Textile & Engineering College, Bhilwara (Raj.)	2023–2027	CGPA: 8.77
Senior Secondary (12th) — S.S.N.V.P. Inter College, Karkauli, Pinahat, Agra (UP)	2022–2023	69%
Secondary (10th) — Manglik Shiksha Kendra, Bamrouli, Fatehabad, Agra (UP)	2020–2021	68.7%

CERTIFICATIONS

- Oracle Certified AI Foundations Associate
- OOSC Conference, IIT Kanpur (Sep 2025)
- Responsive Web Design